

NEURAL NETWORKS AND DEEP LEARNING ICP8

NAME : DASARI BHAVANI BHAGAVATHAMMA

STUDENT ID : 700759340

GITHUB LINK : https://github.com/Bhavani700759340/-NeuralNetworkAndDeepLearning_ICP8/tree/main

VIDEO LINK :

<https://drive.google.com/file/d/1bYmaz1pLvso7HGqkLVYnwjXI4VDF6Njm/view?usp=sharing>

LeNet5 Code and output:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import RMSprop, Adam
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")

(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 6s 0us/step

classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]

y_train = y_train.reshape(-1,)

# Reshape converting 2D to 1D
y_test = y_test.reshape(-1,)
y_train = y_train.reshape(-1,)

# This code normalization
x_train = x_train / 255.0
x_test = x_test / 255.0

x_train.shape

(50000, 32, 32, 3)

from tensorflow.keras import layers, models
lenet = keras.models.Sequential([
    keras.layers.Conv2D(6, kernel_size=5, strides=1, activation='relu', input_shape=(32,32,3), padding='same'), #C1
    keras.layers.AveragePooling2D(), #S1
    keras.layers.Conv2D(16, kernel_size=5, strides=1, activation='relu', padding='valid'), #C2
    keras.layers.AveragePooling2D(), #S2
    keras.layers.Conv2D(120, kernel_size=5, strides=1, activation='relu', padding='valid'), #C3
    keras.layers.Flatten(), #Flatten
    keras.layers.Dense(84, activation='relu'), #F1
    keras.layers.Dense(10, activation='softmax') #Output Layer
])
```

```
lenet.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------------|--------------------|---------|
| ===== | | |
| conv2d (Conv2D) | (None, 32, 32, 6) | 456 |
| average_pooling2d (AveragePooling2D) | (None, 16, 16, 6) | 0 |
| conv2d_1 (Conv2D) | (None, 12, 12, 16) | 2416 |

| | | |
|---|-------------------|-------|
| average_pooling2d_1 (Average Pooling2D) | (None, 6, 6, 16) | 0 |
| conv2d_2 (Conv2D) | (None, 2, 2, 120) | 48120 |
| flatten (Flatten) | (None, 480) | 0 |
| dense (Dense) | (None, 84) | 40404 |
| dense_1 (Dense) | (None, 10) | 850 |

```

=====
Total params: 92,246
Trainable params: 92,246
Non-trainable params: 0

```

```
lenet.compile(optimizer='adam', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
```

```
hist = lenet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
```

```

Epoch 1/100
1563/1563 [=====] - 19s 6ms/step - loss: 1.6343 - accuracy: 0.4051 - val_loss: 1.4179 - val_accuracy: 0.4879
Epoch 2/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.3498 - accuracy: 0.5170 - val_loss: 1.2812 - val_accuracy: 0.5415
Epoch 3/100
1563/1563 [=====] - 8s 5ms/step - loss: 1.2189 - accuracy: 0.5675 - val_loss: 1.2239 - val_accuracy: 0.5587
Epoch 96/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1112 - accuracy: 0.9635 - val_loss: 4.8539 - val_accuracy: 0.5863
Epoch 97/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1035 - accuracy: 0.9671 - val_loss: 5.0162 - val_accuracy: 0.5881
Epoch 98/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1029 - accuracy: 0.9678 - val_loss: 4.9166 - val_accuracy: 0.5819
Epoch 99/100
1563/1563 [=====] - 10s 6ms/step - loss: 0.1059 - accuracy: 0.9658 - val_loss: 5.0741 - val_accuracy: 0.5924
Epoch 100/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1006 - accuracy: 0.9670 - val_loss: 4.9955 - val_accuracy: 0.5902


```


```

# summarize history for accuracy
plt.plot(hist.history['accuracy'])
plt.plot(hist.history['val_accuracy'])
plt.title("Accuracy by LeNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# summarize history for Loss
plt.plot(hist.history['loss'])
plt.plot(hist.history['val_loss'])
plt.title('Loss by LeNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()

```





```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1,)
y_predictions= np.argmax(y_predictions, axis=1)
```


```
confusion_matrix(y_test, y_predictions)
```

```
313/313 [=====] - 1s 2ms/step
```

```
array([[618, 36, 64, 32, 28, 21, 11, 24, 115, 51],
       [ 24, 740, 12, 17, 7, 10, 12, 8, 43, 127],
       [ 69, 21, 495, 76, 77, 110, 53, 55, 21, 23],
       [ 25, 27, 83, 397, 61, 215, 67, 63, 23, 39],
       [ 32, 10, 114, 90, 441, 80, 75, 125, 12, 21],
       [ 18, 8, 75, 200, 39, 523, 38, 70, 13, 16],
       [ 12, 23, 82, 102, 40, 76, 615, 13, 21, 16],
       [ 25, 7, 37, 65, 68, 95, 8, 664, 8, 23],
       [ 77, 55, 11, 22, 12, 16, 14, 9, 737, 47],
       [ 33, 142, 13, 30, 11, 23, 9, 29, 38, 672]])
```

```
# confusion matrix and accuracy
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



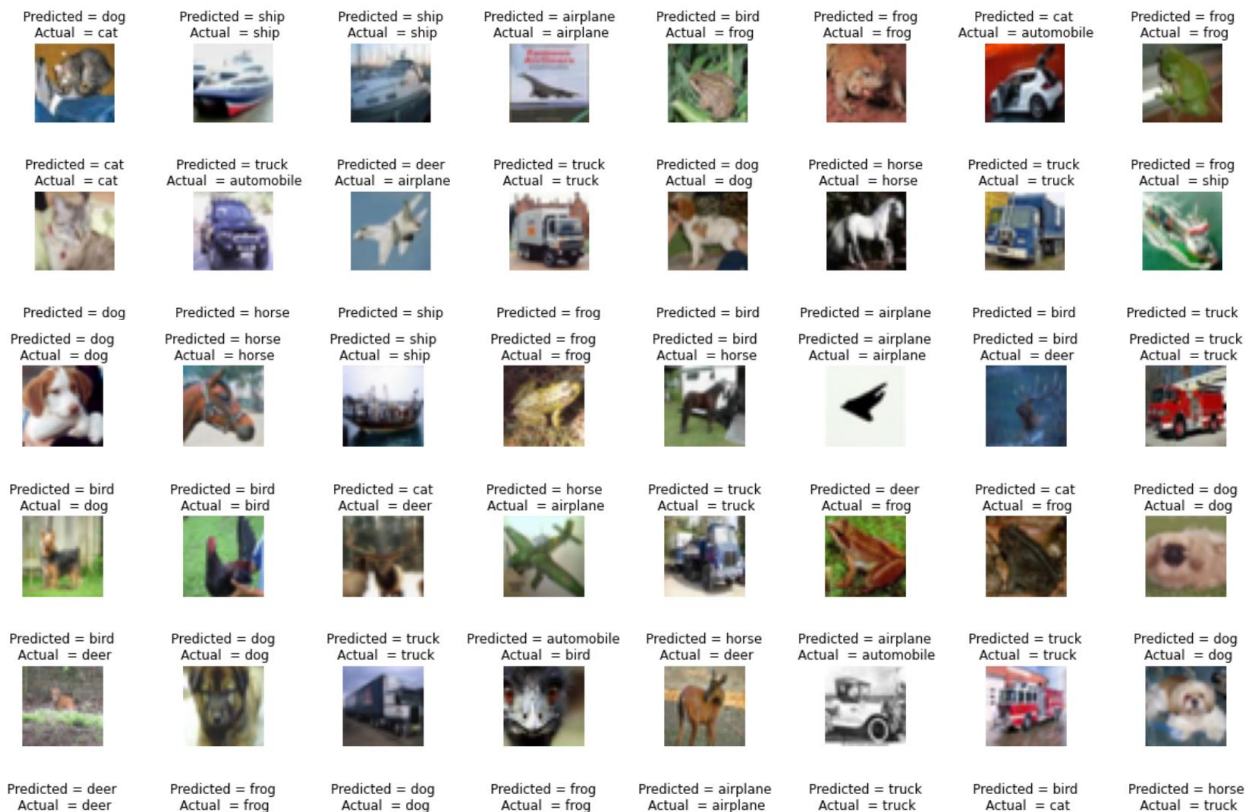
```
print("Test accuracy:", accuracy_score(y_test, y_predictions))
```

Test accuracy: 0.5902

```
L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {} \n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```





AlexNet Code and Output:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.layers.convolutional import Convolution2D
from tensorflow.keras.layers.convolutional import MaxPooling2D
```

```
#Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(3,3),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(10,activation='softmax'))
```

```
AlexNet.compile(optimizer='SGD', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------------|------------------|---------|
| ===== | | |
| conv2d_3 (Conv2D) | (None, 8, 8, 16) | 448 |
| max_pooling2d (MaxPooling2D) | (None, 4, 4, 16) | 0 |
| conv2d_4 (Conv2D) | (None, 4, 4, 60) | 24060 |
| max_pooling2d_1 (MaxPooling2D) | (None, 2, 2, 60) | 0 |
| conv2d_5 (Conv2D) | (None, 2, 2, 60) | 32460 |
| conv2d_6 (Conv2D) | (None, 2, 2, 30) | 16230 |
| conv2d_7 (Conv2D) | (None, 2, 2, 20) | 5420 |
| max_pooling2d_2 (MaxPooling2D) | (None, 1, 1, 20) | 0 |

| | | |
|---------------------|-------------|-------|
| flatten_1 (Flatten) | (None, 20) | 0 |
| dense_2 (Dense) | (None, 200) | 4200 |
| dropout (Dropout) | (None, 200) | 0 |
| dense_3 (Dense) | (None, 200) | 40200 |
| dropout_1 (Dropout) | (None, 200) | 0 |
| dense_4 (Dense) | (None, 10) | 2010 |

```

=====
Total params: 125,028
Trainable params: 125,028
Non-trainable params: 0

```

```
history1 = AlexNet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
```

```

Epoch 1/100
1563/1563 [=====] - 13s 7ms/step - loss: 2.3008 - accuracy: 0.1261 - val_loss: 2.2969 - val_accuracy: 0.1592
Epoch 2/100
1563/1563 [=====] - 10s 7ms/step - loss: 2.2508 - accuracy: 0.1689 - val_loss: 2.0838 - val_accuracy: 0.1888
Epoch 3/100
1563/1563 [=====] - 10s 6ms/step - loss: 2.0630 - accuracy: 0.2030 - val_loss: 1.9858 - val_accuracy: 0.2479
Epoch 4/100
1563/1563 [=====] - 11s 7ms/step - loss: 1.9620 - accuracy: 0.2432 - val_loss: 1.8994 - val_accuracy: 0.2673
Epoch 98/100
1563/1563 [=====] - 11s 7ms/step - loss: 0.2925 - accuracy: 0.8944 - val_loss: 2.4082 - val_accuracy: 0.5349
Epoch 99/100
1563/1563 [=====] - 10s 7ms/step - loss: 0.2918 - accuracy: 0.8943 - val_loss: 2.4604 - val_accuracy: 0.5338
Epoch 100/100
1563/1563 [=====] - 11s 7ms/step - loss: 0.2835 - accuracy: 0.9001 - val_loss: 2.6043 - val_accuracy: 0.5287

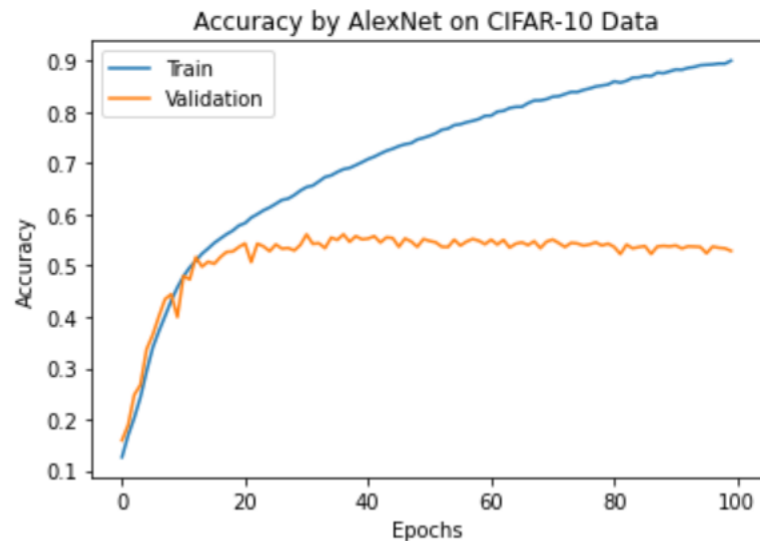
```

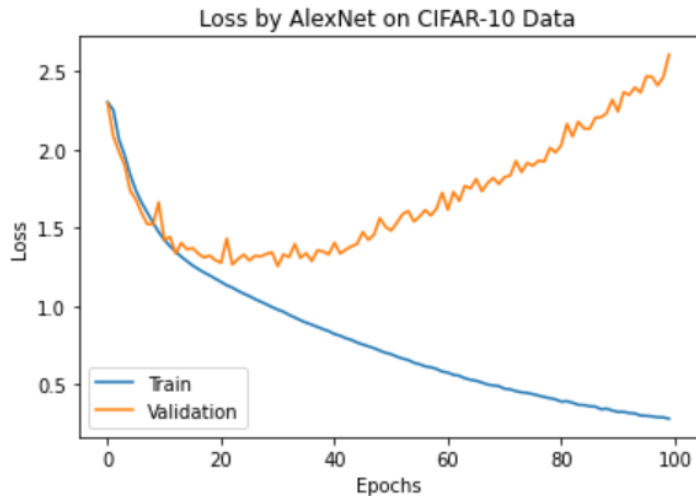
```

# summarize history for accuracy
plt.plot(history1.history['accuracy'])
plt.plot(history1.history['val_accuracy'])
plt.title("Accuracy by AlexNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# summarize history for Loss
plt.plot(history1.history['loss'])
plt.plot(history1.history['val_loss'])
plt.title('Loss by AlexNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()

```



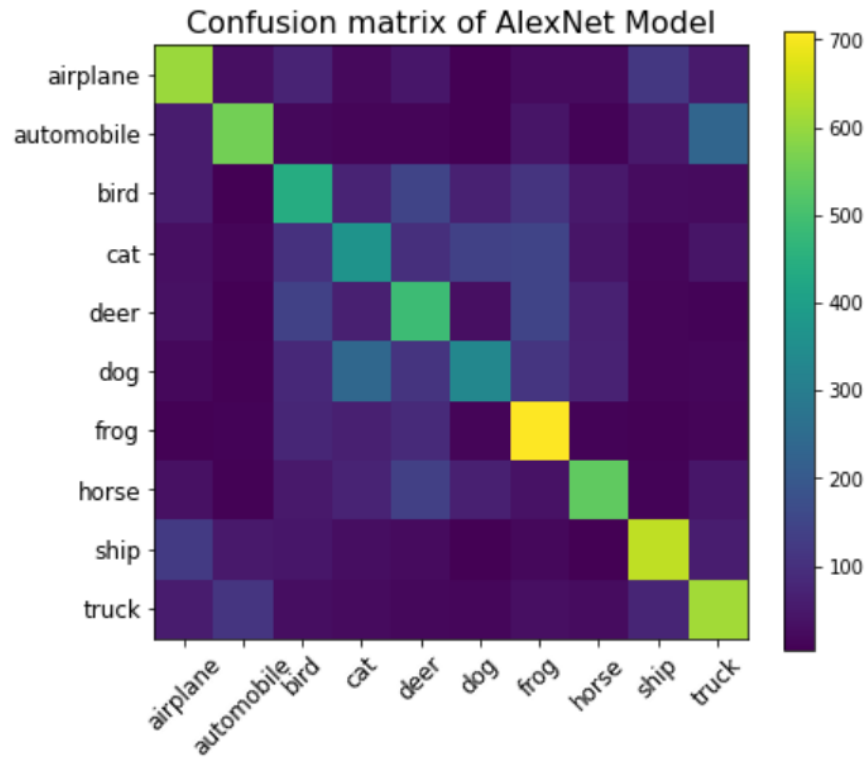


```
y_predictions1 = AlexNet.predict(x_test)
y_predictions1.reshape(-1,)
y_predictions1= np.argmax(y_predictions1, axis=1)

confusion_matrix(y_test, y_predictions1)
```

```
313/313 [=====] - 1s 2ms/step
array([[602, 33, 72, 21, 47, 4, 25, 24, 118, 54],
       [ 56, 560, 19, 13, 13, 4, 42, 11, 50, 232],
       [ 56, 3, 441, 74, 148, 69, 110, 50, 26, 23],
       [ 31, 13, 104, 367, 97, 139, 148, 43, 15, 43],
       [ 34, 4, 139, 64, 488, 32, 149, 68, 12, 10],
       [ 18, 6, 85, 240, 110, 331, 112, 70, 13, 15],
       [ 8, 9, 78, 66, 88, 14, 709, 9, 7, 12],
       [ 34, 7, 51, 74, 138, 65, 38, 538, 9, 46],
       [123, 51, 45, 29, 25, 3, 19, 5, 641, 59],
       [ 58, 111, 28, 24, 19, 16, 33, 26, 75, 610]])
```

```
]: # confusion matrix and accuracy
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix of AlexNet Model', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions1))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```




















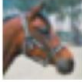


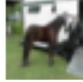



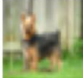

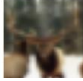














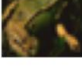

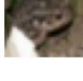


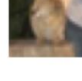
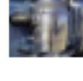



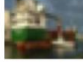





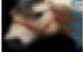


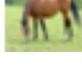
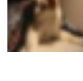
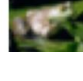

```
print("Test accuracy by AlexNet:", accuracy_score(y_test, y_predictions))
```

Test accuracy by AlexNet: 0.5902

```
L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```


| | | | | | | | |
|--|---|--|--|---|--|--|---|
| Predicted = dog Actual = cat  | Predicted = ship Actual = ship  | Predicted = ship Actual = ship  | Predicted = airplane Actual = airplane  | Predicted = bird Actual = frog  | Predicted = frog Actual = frog  | Predicted = cat Actual = automobile  | Predicted = frog Actual = frog  |
| Predicted = cat Actual = cat  | Predicted = truck Actual = automobile  | Predicted = deer Actual = airplane  | Predicted = truck Actual = truck  | Predicted = dog Actual = dog  | Predicted = horse Actual = horse  | Predicted = truck Actual = truck  | Predicted = frog Actual = ship  |
| Predicted = dog Actual = dog  | Predicted = horse Actual = horse  | Predicted = ship Actual = ship  | Predicted = frog Actual = frog  | Predicted = bird Actual = horse  | Predicted = airplane Actual = airplane  | Predicted = bird Actual = deer  | Predicted = truck Actual = truck  |
| Predicted = bird Actual = dog  | Predicted = bird Actual = bird  | Predicted = cat Actual = deer  | Predicted = horse Actual = airplane  | Predicted = truck Actual = truck  | Predicted = deer Actual = frog  | Predicted = cat Actual = frog  | Predicted = dog Actual = dog  |
| Predicted = bird Actual = deer  | Predicted = dog Actual = dog  | Predicted = truck Actual = truck  | Predicted = automobile Actual = bird  | Predicted = horse Actual = deer  | Predicted = airplane Actual = automobile  | Predicted = truck Actual = truck  | Predicted = dog Actual = dog  |
| Predicted = deer Actual = deer  | Predicted = frog Actual = frog  | Predicted = dog Actual = dog  | Predicted = frog Actual = frog  | Predicted = airplane Actual = airplane  | Predicted = truck Actual = truck  | Predicted = bird Actual = cat  | Predicted = horse Actual = truck  |
| Predicted = horse Actual = horse  | Predicted = frog Actual = frog  | Predicted = truck Actual = truck  | Predicted = ship Actual = ship  | Predicted = dog Actual = airplane  | Predicted = frog Actual = cat  | Predicted = ship Actual = ship  | Predicted = ship Actual = ship  |
| Predicted = dog Actual = horse  | Predicted = ship Actual = horse  | Predicted = dog Actual = deer  | Predicted = cat Actual = frog  | Predicted = horse Actual = horse  | Predicted = bird Actual = cat  | Predicted = automobile Actual = frog  | Predicted = truck Actual = cat  |

Vgg16 Code and Output:

```
import keras
from keras.models import Sequential
from keras.preprocessing import image
from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D,BatchNormalization
from keras.datasets import cifar10
from keras import optimizers
from matplotlib import pyplot as plt
```


```
#generate cifar10 data
(x_train,y_train),(x_test,y_test) = cifar10.load_data()
```

```
#config parameters
num_classes = 10
input_shape = x_train.shape[1:4]
optimizer = optimizers.Adam(lr=0.001)
```

```
#convert label to one-hot
one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```

```
# check data
plt.imshow(x_train[1])
print(x_train[1].shape)
```

(32, 32, 3)



```

# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(2048, activation='relu'))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

```

```
model.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|--------------------|---------|
| ===== | | |
| conv2d_1 (Conv2D) | (None, 32, 32, 64) | 1792 |
| batch_normalization (Batch Normalization) | (None, 32, 32, 64) | 256 |
| conv2d_2 (Conv2D) | (None, 32, 32, 64) | 36928 |
| batch_normalization_1 (Batch Normalization) | (None, 32, 32, 64) | 256 |
| max_pooling2d (MaxPooling2D) | (None, 16, 16, 64) | 0 |
| dropout (Dropout) | (None, 16, 16, 64) | 0 |

```
history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)
```

```
Epoch 1/30
352/352 [=====] - 29s 81ms/step - loss: 1.7083 - accuracy: 0.3346 - val_loss: 2.6400 - val_accuracy: 0.3686
Epoch 2/30
352/352 [=====] - 27s 76ms/step - loss: 1.3159 - accuracy: 0.5254 - val_loss: 1.6935 - val_accuracy: 0.5590
Epoch 27/30
352/352 [=====] - 28s 78ms/step - loss: 0.0962 - accuracy: 0.9703 - val_loss: 0.7319 - val_accuracy: 0.8376
Epoch 28/30
352/352 [=====] - 28s 78ms/step - loss: 0.0871 - accuracy: 0.9728 - val_loss: 0.7403 - val_accuracy: 0.8480
Epoch 29/30
352/352 [=====] - 28s 79ms/step - loss: 0.0811 - accuracy: 0.9749 - val_loss: 0.6592 - val_accuracy: 0.8510
Epoch 30/30
352/352 [=====] - 28s 78ms/step - loss: 0.0753 - accuracy: 0.9770 - val_loss: 0.6631 - val_accuracy: 0.8642
```

```
# evaluate
```

```
print(model.metrics_names)
```

```
model.evaluate(x=x_test,y=one_hot_y_test,batch_size=512)
```

```
['loss', 'accuracy']
```

```
20/20 [=====] - 5s 132ms/step - loss: 0.6627 - accuracy: 0.8592
```

```
[0.6626988053321838, 0.8592000007629395]
```

```
model.save("keras-VGG16-cifar10.h5")
```

```
plt.imshow(x_test[1000])
```

```
result = model.predict(x_test[1000:1001]).tolist()
```

```
predict = 0
```

```
expect = y_test[1000][0]
```

```
for i, _ in enumerate(result[0]):
```


```
    if result[0][i] > result[0][predict]:
```

```
        predict = i
```

```
print("predict class:",predict)
```


```
print("expected class:",expect)
```

```
1/1 [=====] - 1s 740ms/step  
predict class: 5  
expected class: 5
```




```
# save model  
model.save("keras-VGG16-cifar10.h5")
```

```
#plot the training and validation accuracy  
plt.plot(history.history['accuracy'], label='training accuracy')  
plt.plot(history.history['val_accuracy'], label='validation accuracy')  
plt.title('Training and validation accuracy')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()
```



```
#plot the training and validation loss
plt.plot(history.history['loss'], label='training loss')
plt.plot(history.history['val_loss'], label='validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```





```
import numpy as np
from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = y_test.ravel()
cm = confusion_matrix(y_true, y_pred_classes)

# plot the confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(num_classes)
plt.xticks(tick_marks, range(num_classes))
plt.yticks(tick_marks, range(num_classes))
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.show()

# plot a histogram of the predicted probabilities for a sample image
plt.hist(y_pred[1000])
plt.title('Predicted probabilities')
plt.xlabel('Probability')
plt.ylabel('Frequency')
plt.show()
```


313/313 [=====] - 3s 8ms/step



VGG19Model Code and Output:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf

from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization

%matplotlib inline
```


Extract data and train and test dataset


```
cifar100 = tf.keras.datasets.cifar100
(X_train,Y_train) , (X_test,Y_test) = cifar100.load_data()
```

```
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Let's look into the dataset images

```
plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,1+i)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```





Training , Validating and Splitting trained and tested data

```
from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train,Y_train,test_size=0.2)
```

```
from keras.utils.np_utils import to_categorical
y_train = to_categorical(y_train, num_classes = 10)
y_val = to_categorical(y_val, num_classes = 10)
```

```
print(x_train.shape)
print(y_train.shape)
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(40000, 32, 32, 3)
(40000, 10)
(10000, 32, 32, 3)
(10000, 10)
(10000, 32, 32, 3)
(10000, 1)
```

```

train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)

val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)

```

```

from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)

```

We have used only 16 layers out of 19 layers in the CNN

```

vgg_model = tf.keras.applications.VGG19(
    include_top=False,
    weights="imagenet",
    input_shape=(32,32,3),
)

vgg_model.summary()

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80142336/80134624 [=====] - 1s 0us/step
Model: "vgg19"

| Layer (type) | Output Shape | Param # |
|----------------------------|---------------------|---------|
| ===== | | |
| input_1 (InputLayer) | [(None, 32, 32, 3)] | 0 |
| ----- | | |
| block1_conv1 (Conv2D) | (None, 32, 32, 64) | 1792 |
| block1_conv2 (Conv2D) | (None, 32, 32, 64) | 36928 |
| ----- | | |
| block1_pool (MaxPooling2D) | (None, 16, 16, 64) | 0 |
| ----- | | |
| block2_conv1 (Conv2D) | (None, 16, 16, 128) | 73856 |
| block2_conv2 (Conv2D) | (None, 16, 16, 128) | 147584 |
| ----- | | |
| block2_pool (MaxPooling2D) | (None, 8, 8, 128) | 0 |
| ----- | | |
| block3_conv1 (Conv2D) | (None, 8, 8, 256) | 295168 |

```

model = tf.keras.Sequential()
model.add(vgg_model)
model.add(Flatten())
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(256, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))

model.summary()

```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------|-------------------|----------|
| vgg19 (Functional) | (None, 1, 1, 512) | 20024384 |
| flatten (Flatten) | (None, 512) | 0 |
| dense (Dense) | (None, 1024) | 525312 |
| dense_1 (Dense) | (None, 1024) | 1049600 |
| dense_2 (Dense) | (None, 256) | 262400 |
| dense_3 (Dense) | (None, 10) | 2570 |

=====
 Total params: 21,864,266
 Trainable params: 21,864,266
 Non-trainable params: 0

```

optimizer = tf.keras.optimizers.SGD(lr = 0.001, momentum = 0.9)
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

```

```

history = model.fit(
    train_datagen.flow(x_train, y_train, batch_size = 128),
    validation_data = val_datagen.flow(x_val, y_val, batch_size = 128),
    epochs = 25,
    verbose = 1,
    callbacks = [learning_rate_reduction]
)

Epoch 1/25
313/313 [=====] - 84s 165ms/step - loss: 1.6671 - accuracy: 0.3875 - val_loss: 1.1241 - val_accuracy: 0.6056
Epoch 2/25
313/313 [=====] - 50s 159ms/step - loss: 0.9631 - accuracy: 0.6640 - val_loss: 0.7466 - val_accuracy: 0.7404
Epoch 3/25
313/313 [=====] - 50s 160ms/step - loss: 0.7464 - accuracy: 0.7430 - val_loss: 0.6792 - val_accuracy: 0.7716
Epoch 24/25
313/313 [=====] - 50s 159ms/step - loss: 0.1964 - accuracy: 0.9304 - val_loss: 0.4461 - val_accuracy: 0.8724
Epoch 25/25
313/313 [=====] - 50s 159ms/step - loss: 0.1879 - accuracy: 0.9337 - val_loss: 0.4827 - val_accuracy: 0.8671

```


```

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Accuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()

```

<matplotlib.legend.Legend at 0x7fc5601fe610>




```

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure()
plt.plot(loss,color = 'green',label = 'Training Loss')
plt.plot(val_loss,color = 'red',label = 'Validation Loss')
plt.legend()

```

<matplotlib.legend.Legend at 0x7fc5bd878ad0>



```

x_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = model.predict_classes(x_test)
y_pred[:10]

```

```
array([3, 8, 8, 0, 6, 6, 1, 6, 3, 1])
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))
```

```
Testing Accuracy : 0.8591
```

```
cm = confusion_matrix(Y_test, y_pred)
cm
```

```
array([[882, 11, 20, 1, 12, 1, 5, 9, 43, 16],
       [ 5, 944, 0, 2, 0, 1, 2, 0, 8, 38],
       [22, 3, 800, 11, 54, 30, 56, 9, 10, 5],
       [10, 7, 42, 536, 57, 228, 74, 26, 6, 14],
       [ 6, 1, 22, 5, 884, 20, 35, 23, 2, 2],
       [ 5, 4, 9, 54, 36, 845, 19, 26, 0, 2],
       [ 5, 1, 11, 7, 16, 10, 941, 3, 3, 3],
       [ 4, 1, 12, 7, 31, 28, 8, 902, 2, 5],
       [25, 16, 1, 1, 4, 0, 4, 1, 929, 19],
       [ 6, 41, 1, 2, 4, 0, 6, 2, 10, 928]])
```

```
import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
```

```
horizontalalignment="center",  
color="white" if cm[i, j] > thresh else "black")
```

```
plt.tight_layout()  
plt.ylabel('True label')  
plt.xlabel('Predicted label')
```

```
plt.figure(figsize=(8,8))  
plot_confusion_matrix(cm,classes)
```

Confusion matrix, without normalization

